

Seminararbeit

Dynamische Programmanalysen für nebenläufige Programme - Data
Race Prediction mit TSan




Frank Ling

June 13, 2023


- ① Introduction
- ② Background
- ③ FastTrack and TSan V2
 - Limitations
- ④ Conclusion

- What are data races?




Source: <https://programming.guide/go/data-races-explained.html>

- What are data races?
- Why fix data races?



Source: <https://programming.guide/go/data-races-explained.html>

- What are data races?
- Why fix data races?
- How to detect data races?



Source: <https://programming.guide/go/data-races-explained.html>

```
1 int x;
2 pthread_mutex_t y;
3
4 void *Thread1(void *x) {
5     x++;
6     pthread_mutex_lock(&y);
7     pthread_mutex_unlock(&y);
8     return NULL;
9 }
10
11 void *Thread2(void *x) {
12     pthread_mutex_lock(&y);
13     x--;
14     pthread_mutex_unlock(&y);
15     return NULL;
16 }
```

Listing: program exhibiting a data race

```
1 int x;
2 pthread_mutex_t y;
3
4 void *Thread1(void *x) {
5     x++;
6     pthread_mutex_lock(&y);
7     pthread_mutex_unlock(&y);
8     return NULL;
9 }
10
11 void *Thread2(void *x) {
12     pthread_mutex_lock(&y);
13     x--;
14     pthread_mutex_unlock(&y);
15     return NULL;
16 }
```

	1#	2#
1.	w(x)	
2.	acq(y)	
3.	rel(y)	
4.		acq(y)
5.		w(x)
6.		rel(y)

Table: obtained trace

Listing: program exhibiting a data race

	1#	2#
4.		acq(y)
5.		w(x)
1.	w(x)	
6.		rel(y)
2.	acq(y)	
3.	rel(y)	

Table: Trace 1 reordered

Data race! \Leftarrow

	1#	2#
4.		acq(y)
5.		w(x)
1.	w(x)	
6.		rel(y)
2.	acq(y)	
3.	rel(y)	

Table: Trace 1 reordered

- Dynamic data race prediction

- Dynamic data race prediction
- Vector clocks

- Dynamic data race prediction
- Vector clocks
- Epochs

- Dynamic data race prediction
- Vector clocks
- Epochs
- Lamport's HB relation

Program order condition

- For events e and f in same thread, if e appears before f in trace then $e <_{HB} f$
- $\text{inc}(V, j)$ applied after processing event in same thread

Critical section

- For $\text{rel}(y)$ in thread i and $\text{acq}(y)$ in thread j if $\text{rel}(y)$ appears before $\text{acq}(y)$ then $\text{rel}(y) <_{HB} \text{acq}(y)$
- $\text{sync}(V_1, V_2)$ is applied after $\text{rel}(y)$ in thread

FastTrack

- Lamport's HB relation
- epoch-based
- semi-adaptive
 - ▶ dynamically sized epochs expensive
 - ▶ initially epochs only
 - ▶ read \rightarrow VC, stays VC after (subsequent concurrent reads frequent)
 - ▶ writes are epochs only (all writes are totally ordered)

ThreadSanitizer (TSan) V2

- slightly modified version of FastTrack
- shadow memory

- Shadow word (64 bits)

TID (Thread Id)	16 bits
Scalar Clock	42 bits
IsWrite	1 bit
Access Size	2 bits
Address Offset	3 bits

- Shadow word (64 bits)

TID (Thread Id)	16 bits
Scalar Clock	42 bits
IsWrite	1 bit
Access Size	2 bits
Address Offset	3 bits

- N** shadows words



application word (memory location)

- Shadow word (64 bits)

TID (Thread Id)	16 bits
Scalar Clock	42 bits
IsWrite	1 bit
Access Size	2 bits
Address Offset	3 bits

- N** shadows words



application word (memory location)

- N is configurable (2, 4, 8) but default is 4



For every memory location 1 write and up to concurrent 3 reads are stored

- Shadow word (64 bits)

TID (Thread Id)	16 bits
Scalar Clock	42 bits
IsWrite	1 bit
Access Size	2 bits
Address Offset	3 bits

- N** shadows words



application word (memory location)

- N is configurable (2, 4, 8) but default is 4



For every memory location 1 write and up to concurrent 3 reads are stored

- for every subsequent concurrent read a random read in shadow memory will be evicted

- Unsound due to epochs

```
1 void thread1() {  
2     y := x+5;  
3 }  
4  
5 void thread2() {  
6     if (y == 5)  
7         x := 10;  
8     else  
9         while(true);  
10 }
```

Listing: program showing unordered writes

	1#	2#
1.	r(x)	
2.	w(y)	
3.		r(y)
4.		w(x)

Table: Obtained trace

- Incomplete due to:
 - ▶ HB relation (see first example)
 - ▶ shadow memory

- no race detection tool right now is entirely complete and sound
- FastTrack and TSan produce good results even though missing data races under certain circumstances
- they are reasonably efficient compared to other data race detection tools