

Dynamische Programmanalysen für nebenläufige Programme - Data Race Prediction mit TSan V2

Seminararbeit

Student:	Frank Ling
Matrikelnummer:	79496
Universität:	Hochschule Karlsruhe – Technik und Wirtschaft
Studiengang:	Informatik, Master
Semester:	Sommersemester 2023
Betreuer:	Prof. Martin Sulzmann
Bearbeitet am:	June 11, 2023

Contents

1	Introduction	1
2	Motivation and Examples	2
3	Background	4
3.1	Dynamic Data Race Prediction	4
3.2	Happens-Before Relation and Data Race	4
3.3	Vector-Clock	5
3.4	Epoch	6
4	FastTrack + TSan	7
5	Conclusion	12
	List of Literature	13
	List of Figures	14
	List of Tables	15
	List of Listings	16

1 Introduction

Nowadays concurrent programs are very common in order to make use of 'hyper-threading and multi-core architectures' [1, p. 14]. 'Due to the highly non-deterministic behavior of concurrent programs' [2, p. 1] data races may arise but can also be hard to find, as they also 'may only arise under a specific schedule' [2, p. 1]. This will be elaborated in the following section. We specifically look at the dynamic data race prediction algorithms of FastTrack and ThreadSanitizer (TSan) V2. TSan V2 uses a slightly modified version of the FastTrack algorithm and differentiates itself from the first version by utilizing happens-before methods instead of the lockset method.

Firstly the motivation and background for dynamic data race prediction tools will be shown. Then the FastTrack and TSan algorithms will be demonstrated along with examples showing the limitations of FastTrack and TSan. In the following the same notation as in [2] will be used for traces and events.

2 Motivation and Examples

As stated in the introduction concurrent programs are commonly used and are inherently prone to data races. The problem with data races is that they can cause undefined behavior and crashes while also being hard to reproduce. This makes data races hard to find and they might not always be apparent. This chapter will show a few examples from [3] but rewritten in the programming language C++. The following example shows a program, which exhibits a data race (see chapter 3):

```
1  #include "pthread.h"
2
3  int var1;
4  pthread_mutex_t var2;
5
6  void *Thread1(void *x) {
7      // w(x)1
8      var1++;
9      // acq(y)2
10     pthread_mutex_lock(&var2);
11     // rel(y)3
12     pthread_mutex_unlock(&var2);
13     return NULL;
14 }
15
16 void *Thread2(void *x) {
17     // acq(y)4
18     pthread_mutex_lock(&var2);
19     // w(x)5
20     var1--;
21     // rel(y)6
22     pthread_mutex_unlock(&var2);
23     return NULL;
24 }
25
26 int main() {
27     pthread_t t[2];
28     pthread_create(&t[0], NULL, Thread1, NULL);
29     pthread_create(&t[1], NULL, Thread2, NULL);
30     pthread_join(t[0], NULL);
31     pthread_join(t[1], NULL);
32 }
```

Listing 1: Program containing conflicting events

Generally the corresponding event in the resulting trace is seen above each operation as a comment in the code. For example acquiring and releasing a lock is accomplished by the

`pthread_mutex_lock()`- and the `pthread_mutex_unlock()`-method, which is used in line 10, where a lock on the variable `var2` is acquired. The corresponding event in the trace would be $1\#acq(y)_2$ which is indicated by the method name and the comment above in line 9. The method name shows which thread the event is from and the comment specifies the event. As a write event a variable can simply be incremented or assigned a value. A simple read event could be printing a variable in the console as it accesses the value of the variable. The following trace can be generated for the program 1:

	1#	2#
1.	w(x)	
2.	acq(y)	
3.	rel(y)	
4.		acq(y)
5.		w(x)
6.		rel(y)

Table 1: Possible trace of listing 1

In the trace 1 the events $1\#w(x)_1$ and $2\#w(x)_5$ do not appear right next to each other which might indicate that these two events are not part of a data race. By rerunning the program 1 eventually the following trace can be obtained:

	1#	2#
4.		acq(y)
5.		w(x)
1.	w(x)	
6.		rel(y)
2.	acq(y)	
3.	rel(y)	

Table 2: Trace 1 reordered

The trace 2 shows the two events $1\#w(x)_1$ and $2\#w(x)_5$ are now next to each other in the trace, thus they are part of a data race. This also shows that there are data races that can only be detected under a specific schedule. Data race prediction tools use various methods to try to predict such cases.

3 Background

3.1 Dynamic Data Race Prediction

Race prediction can be done statically or dynamically. Static race prediction uses static analysis, meaning it tries to find data races before running the code. For more details on static race prediction see [4].

Dynamic data race prediction tools run a program and then, based on the obtained trace, try to predict if there may be a reordered trace that exhibits data races, see [3, p. 2].

The approach of finding all possible reorderings is used in exhaustive predictive methods. Since modern programs may utilize multiple threads each executing large, computationally intensive blocks of code, finding all possible reorderings might not always be feasible. Therefore effective predictive methods are employed which compromise *completeness* and *soundness* for better coverage. FastTrack and TSan are examples for dynamic data race prediction using effective predictive methods.

Sulzmann and Stadtmüller [2, p. 2] define completeness and soundness as follows:

Sound means that races reported by the algorithm can be observed via some appropriate reordering of the trace. If unsound, we refer to wrongly a classified race as a *false positive*. Complete means that all valid reorderings that exhibit some race can be predicted by the algorithm. If incomplete, we refer to any not reported race as a *false negative*.

3.2 Happens-Before Relation and Data Race

As stated in [3, p. 4], the happens-before relation is used to partially order the events in a system. The happens-before relation is defined as a strict partial order so the ordering is transitive but irreflexive.

Events can either be ordered and therefore a happens-before relation can be established between them, or they can not be ordered in which case both events 'are said to be *concurrent*' [5, p. 2]. Further two events '[...] *conflict* if they both access (read or write) the same variable, and at least one of the operations is a write. Using this terminology, a trace has a race condition if it has two concurrent conflicting accesses.' [6, p. 2]. An example of a data race can be seen in the trace 2.

In FastTrack and TSan Lamport's happens-before relation [5], referred to as HB relation, is used. The HB relation additionally to the above imposes the program order condition and critical section order condition. The program order condition defines that two events e and f belonging to the same thread, with e appearing before f in the trace, then $e <^{HB} f$ applies. $<^{HB}$ denotes the HB ordering relation. The critical section order condition imposes that for two acquire and release events $acq(y)$ and $rel(y)$ on the same lock y where both events result from different threads with $rel(y)$ appearing before $acq(y)$ in the trace, then $rel(y) <^{HB} acq(y)$ applies [3, p. 4].

3.3 Vector-Clock

As stated by Flanagan and Freund [6, p. 1]:

Typically, the happens-before relation is represented using vector clocks (VCs) [...]. Vector clocks are expensive, however, because they record information about each thread in a system. Thus, if the target program has n threads, then each VC requires $O(n)$ storage space and each VC operation requires $O(n)$ time. Motivated in part by the performance limitations of vector clocks, a variety of alternative imprecise race detectors have been developed, which may provide better coverage but can report false alarms on race-free programs.

Sulzmann and Stadtmüller define vector clocks as follows [2, p. 7]:

A vector clock V is a list of time stamps of the following form.

$$v ::= [i_1, \dots, i_n] \quad (1)$$

We assume vector clocks are of a fixed size n . Time stamps are natural numbers and each time stamp position j corresponds to the thread with identifier j . We define

$$[i_1, \dots, i_n] \sqcup [j_1, \dots, j_n] = [\max(i_1, j_1), \dots, \max(i_n, j_n)] \quad (2)$$

to synchronize two vector clocks by building the point-wise maximum. We write $V[j]$ to access the time stamp at position j . We write $\text{inc}(V, j)$ as a short-hand for incrementing the vector clock V at position j by one. We define vector clock V_1 to be smaller than vector clock V_2 , written $V_1 < V_2$, if (1) for each thread i , i 's time stamp in V_1 is smaller or equal compared to i 's time stamp in V_2 , and (2) there exists a thread i where i 's time stamp in V_1 is strictly smaller compared to i 's time stamp in V_2 . [...] For each thread i we maintain a vector clock $Th(i)$. [...] Initially, for each vector clock $Th(i)$ all time stamps are set to 0 but position i where the time stamp is set to 1.

Sulzmann and Stadtmüller further state that for two events e and f with vector clocks V_1 assigned to e and V_2 assigned to f , if $V_1 < V_2$ then it can be said that $e <^{HB} f$. Further for synchronization the following applies. $V_1 = V_2 \sqcup V_3$ then $V_1 \leq V_2$ and $V_1 \leq V_3$.

For the cases where neither $V_1 < V_2$ nor $V_2 < V_1$ is true it can be said that V_1 and V_2 are incomparable [3, p. 8]. The events assigned with V_1 and V_2 are therefore concurrent. If at least one of the events is a write event, then a data race has been found.

With Lamport's HB relation for events in the same thread j with the respective vector clock V_n , the operation $\text{inc}(V_n, j)$ is applied after processing the event. This is to enforce the program order condition. That results in events having a pre vector clock $\text{pre_}V$ and a post vector clock $\text{post_}V$. As an example in trace 1, $1\#w(x)$ will have the pre vector clock $[1, 0]$, as we are in the first thread we must initialize set the time stamp at position one to

zero. Then $\text{inc}(\text{pre_}V_1, 1)$ is applied so the post vector clock would be $\text{post_}V_1 = [2, 0]$. For following events in the same thread the pre vector clock is the same as the post vector clock of the previous event. So for $2\#\text{acq}(y)$ the pre vector clock would be $[2, 0]$ and the post vector clock $[3, 0]$. To enforce the critical section order condition the pre vector clock of an acquire event following a release event on the same lock has to be synchronized. The post vector clock of the acquire event will be the incremented synchronized vector clock. In trace 1 the pre vector clock of $3\#\text{rel}(y)$ is $[3, 0]$ and $[0, 1]$ for $4\#\text{acq}(y)$. So the synchronized vector clock is $V_2 = [3, 1]$ and after applying $\text{inc}(V_2, 2)$, the vector clock for $4\#\text{acq}(y)$ would be $[3, 2]$. For further details on vector clocks, see [3, p. 8].

3.4 Epoch

4 FastTrack + TSan

```
1 #include <pthread.h>
2
3 int var;
4
5 void *Thread1(void *x) {
6     // w(x)1
7     var++;
8     return NULL;
9 }
10
11 void *Thread2(void *x) {
12     // w(x)2
13     var--;
14     // w(x)3
15     // this write will not be reported as part of a data race
16     var++;
17     return NULL;
18 }
19
20 int main() {
21     pthread_t t[2];
22     pthread_create(&t[0], NULL, Thread1, NULL);
23     pthread_create(&t[1], NULL, Thread2, NULL);
24     pthread_join(t[0], NULL);
25     pthread_join(t[1], NULL);
26 }
```

Listing 2: Program demonstrating the incompleteness of TSan

As seen in [7] Lamport’s happens-before relation is also unsound. Because FastTrack and TSan both use the HB relation they might produce false positives as a result. The following program 3 shows an example in which TSan produces a false negative:

```
1 #include <pthread.h>
2
3 int var1;
4 int var2;
5
6 void *Thread1(void *x) {
7     // r(x)1
8     // w(y)2
9     var2 = var1 + 5;
10    return NULL;
11 }
12
```

```

13 void *Thread2(void *x) {
14     // r(y)3
15     if (var2 == 5)
16         // w(x)4
17         var1 = 10;
18     else
19         while (true);
20     return NULL;
21 }
22
23 int main() {
24     pthread_t t[2];
25     pthread_create(&t[0], NULL, Thread1, NULL);
26     pthread_create(&t[1], NULL, Thread2, NULL);
27     pthread_join(t[0], NULL);
28     pthread_join(t[1], NULL);
29 }

```

Listing 3: Program demonstrating the unsoundness of TSan

The program 3 may produce the following trace 3:

	1#	2#
1.	r(x)	
2.	w(y)	
3.		r(y)
4.		w(x)

Table 3: Possible trace of listing 3

Describe shadow cells and shadow words, then example...

```

1 #include "pthread.h"
2 #include "iostream"
3 using namespace std;
4
5 int var1;
6 int threadcount = 6;
7 pthread_mutex_t m;
8
9 void *Thread1(void *x) {
10     // w(x)
11     var1=1;
12     // acq(y)
13     pthread_mutex_lock(&m);
14     // rel(y)

```

```

15     pthread_mutex_unlock(&m);
16     return NULL;
17 }
18
19 void *ThreadN(void *x) {
20     // acq(y)
21     pthread_mutex_lock(&m);
22     // rel(y)
23     pthread_mutex_unlock(&m);
24     // r(x)
25     cout<<var1;
26     return NULL;
27 }
28
29 void *Thread6(void *x) {
30     // acq(y)
31     pthread_mutex_lock(&m);
32     // rel(y)
33     pthread_mutex_unlock(&m);
34     // w(x)
35     var1=1;
36     return NULL;
37 }
38
39 int main() {
40     pthread_t t[threadcount];
41     for (int i = 0; i < threadcount; i++) {
42         if (i == 0) {
43             pthread_create(&t[i], NULL, Thread1, NULL);
44         } else if (i == threadcount) {
45             pthread_create(&t[i], NULL, Thread6, NULL);
46         } else {
47             pthread_create(&t[i], NULL, ThreadN, NULL);
48         }
49     }
50
51     for (pthread_t& thread : t) {
52         pthread_join(thread, NULL);
53     }
54 }

```

Listing 4: Program showing a missed data race due to the limit on shadow words

ThreadSanitizer fails to show stack trace for data races that are at least 16 threads apart from each other

```

1 #include "pthread.h"

```

```

2 #include "iostream"
3 using namespace std;
4
5 int var1;
6 int threadCount = 18;
7 pthread_mutex_t m;
8
9 void *Thread1(void *x) {
10     var1++;
11     pthread_mutex_lock(&m);
12     pthread_mutex_unlock(&m);
13     return NULL;
14 }
15
16 void *ThreadN(void *x) {
17     pthread_mutex_lock(&m);
18     pthread_mutex_unlock(&m);
19     cout<<var1;
20     return NULL;
21 }
22
23 void *Thread18(void *x) {
24     var1--;
25     return NULL;
26 }
27
28 int main() {
29     pthread_t t[threadCount];
30     for (int i = 0; i < threadCount; i++) {
31         if (i == 0) {
32             pthread_create(&t[i], NULL, Thread1, NULL);
33         } else if (i == threadCount - 1) {
34             pthread_create(&t[i], NULL, Thread18, NULL);
35         } else {
36             pthread_create(&t[i], NULL, ThreadN, NULL);
37         }
38     }
39
40     for (pthread_t& thread : t) {
41         pthread_join(thread, NULL);
42     }
43 }

```

Listing 5: Program showing thread count limitation

5 Conclusion

List of Literature

- [1] A. R. Molla, G. Sharma, P. Kumar, and S. Rawat, Eds., *Distributed Computing and Intelligent Technology : 19th International Conference, ICDCIT 2023, Bhubaneswar, India, January 18–22, 2023, Proceedings*, Cham, 2023. [Online]. Available: <https://link.springer.com/book/10.1007/978-3-031-24848-1>.
- [2] M. Sulzmann and K. Stadtmüller, “Efficient, Near Complete and Often Sound Hybrid Dynamic Data Race Prediction (extended version),” *CoRR*, vol. abs/2004.06969, 2020. arXiv: 2004.06969. [Online]. Available: <https://arxiv.org/abs/2004.06969>.
- [3] M. Sulzmann, *Dynamic data race prediction*. [Online]. Available: <https://sulzmann.github.io/AutonomieSysteme/lec-data-race.html>, accessed Jun. 5, 2023.
- [4] M. Naik, A. Aiken, and J. Whaley, “Effective static race detection for java,” *SIGPLAN Not.*, vol. 41, no. 6, 308–319, 2006, ISSN: 0362-1340. DOI: 10.1145/1133255.1134018. [Online]. Available: <https://doi.org/10.1145/1133255.1134018>.
- [5] L. Lamport, “Time, Clocks, and the Ordering of Events in a Distributed System,” *Commun. ACM*, vol. 21, no. 7, 558–565, 1978, ISSN: 0001-0782. DOI: 10.1145/359545.359563. [Online]. Available: <https://doi.org/10.1145/359545.359563>.
- [6] C. Flanagan and S. Freund, “FastTrack: Efficient and Precise Dynamic Race Detection,” vol. 53, Jun. 2009, pp. 121–133. DOI: 10.1145/1542476.1542490.
- [7] U. Mathur, D. Kini, and M. Viswanathan, “What Happens-after the First Race? Enhancing the Predictive Power of Happens-before Based Dynamic Race Detection,” *Proc. ACM Program. Lang.*, vol. 2, no. OOPSLA, 2018. DOI: 10.1145/3276515. [Online]. Available: <https://doi.org/10.1145/3276515>.

List of Figures

List of Tables

1	Possible trace of listing 1	3
2	Trace 1 reordered	3
3	Possible trace of listing 3	8

List of Listings

1	Program containing conflicting events	2
2	Program demonstrating the incompleteness of TSan	7
3	Program demonstrating the unsoundness of TSan	7
4	Program showing a missed data race due to the limit on shadow words . .	8
5	Program showing thread count limitation	9