

concurrent programs highly non-deterministic

data race = 2 conflicting events right next to each other in trace

conflicting events = 2 read/write events,  
at least one event is write event

you can find potential data races by re-running the program and obtain different traces

alternative is finding valid reorderings of a trace

Trace A may not contain data races but its reordered Trace B or C may

Possible to predict trace orderings that exhibit data races  $\rightarrow$  dynamic data race prediction

Exhaustive predictive methods  $\rightarrow$  identify as many reorderings as possible

Efficient predictive methods  $\rightarrow O(n)$  runtime  
efficient favourable over being exhaustive

$\rightarrow$  compromise completeness and soundness

complete = all valid reorderings that exhibit data races  
can be predicted  
any not predicted race  $\rightarrow$  false negative

sound = races reported can be observed via appropriate  
reordering of trace, unsound  $\rightarrow$  wrongly classified  
data races  $\rightarrow$  false positive

Lockset method unsound, happens-before method  
incomplete

events processed in stream-based fashion = online  
offline may get better results if trace in its entire form  
is present

Happens-before method

1. Trace from specific program run
2. Derive happens-before ordering relation from trace
3. Two conflicting events are unordered  $\rightarrow$  data race

$e < f \rightarrow e$  ordered before  $f$  =  $e$  connected to  $f$   
 $\rightarrow e$  must be executed before  $f$

Req:

- Strict partial order

$\rightarrow$  partial: not all events need to be ordered

$\rightarrow$  strict partial order: ordering relation a) transitive + b) not reflexive

a)  $e < f, f < g \rightarrow e < g$       b) event cannot happen itself

If two conflicting events,  $e, f$ : not  $e < f$  nor  $f < e$

→  $(e, f)$  = data race pair

→ trace can be reordered such that  $e, f$  appear next to each other

If  $(e, f)$  is data race then  $(f, e)$  is also data race

$(e, f)$  and  $(f, e)$  = distinct representative for same data race

Lamport's happens-before relation

Program order condition

$e, f$  events in same thread where  $e$  appears before  $f$ , then  $e < f$

Critical section order condition

$acq(y), rel(y)$  same lock  $y$ , both events result from diff. threads and  $rel(y)$  appears before  $acq(y)$  in trace, then  $rel(y) < acq(y)$

# Vector Clocks

every thread time stamp in array

Initially all entries 0ed except Thread  $t$ , set to 1

$inc([h_1, \dots, h_{i-1}, h_i, h_{i+1}, \dots, h_n], i) = [h_1, \dots, h_{i-1}, h_i+1, \dots, h_n]$

$sync([i_1, \dots, i_n], [j_1, \dots, j_n]) = [\max(i_1, j_1), \dots, \max(i_n, j_n)]$

program order cond.  $\rightarrow$  inc

critical section order cond.  $\rightarrow$  sync for  $rel(y) < acq(y)$   
then inc

$[pre_1, pre_2] - fun(x) - [post_1, post_2]$

find HB data races by comparing pre vector clocks

$V_1 < V_2$  if  $\forall$  array position  $i$ ,  $V_1[i] \leq V_2[i]$   
and array pos  $j$  exists with  $V_1[j] < V_2[j]$

If neither  $V_1 < V_2$ , nor  $V_1 > V_2$  then  $V_1, V_2$   
are incomparable

Vector clocks offline

1st pass: store all pre vector clocks of read/write events

2nd pass: find reorderings of trace with data races

$\rightarrow$  extra space needed for storing pre vector clocks

$\rightarrow$  extra time needed because two passes

but all HB data races can be identified

Online

- pass through trace and store every read/write event if they are incomparable (data race)
- but only keep event if it represents data race, non-data-race events are only stored until next read/write-event

→ invariant:  $V_1, V_2 \in W(x) \cup R(x)$ ,  $\neg \text{happensBefore}(V_1, V_2)$   
and  $\neg \text{happensBefore}(V_2, V_1)$

→  $\forall (V_1, V_2)$  where  $V_1, V_2 \in W(x)$  write-write-  
data race

→  $\forall (V_1, V_2)$  where  $V_1 \in W(x)$  and  $V_2 \in R(x)$   
write-read data race

